

GROUNDTRUTH

Operations Intelligence for Qadira Developments, on Salesforce

Forward Deployed Consultant — Take-Home Submission

AISTRA Talent

Prepared by Charith Wickramasinghe · September 2026

Submission pack: this document, plus the working prototype (qadira_prototype.html)

CONTENTS

Table of Contents

00 Executive Summary 3

01 The Offering 3

02 The Offering, Made Real 5

03 Pricing..... 6

04 The Business Case for Qadira..... 7

05 How We Sell It..... 10

06 Judgement Calls, Recorded 12

07 Anticipated Questions 12

SECTION 00

Executive Summary

Layla wants to be told things, not asked to go and look, without her team entering more data. Sandeep needs whatever we propose to avoid the shape of the AI add-on the CFO already regrets buying, and to give Northbay a reason to stay in the deal. The CFO isn't in either note, and I've treated that as the harder constraint, not the easier one — the numbers in Section 4 are built to stand up without him having said anything at all.

The offering: **Groundtruth** is a weekly, push-delivered intelligence layer that reads Qadira's existing Salesforce data — no new fields, no new steps for the 22-person sales team — and produces four things Layla asked for by name: a reality-checked pipeline forecast, a ranked list of deals losing momentum now, a broker performance ranking built on incremental lift rather than volume, and a unit-by-unit velocity map. It sits outside Salesforce's UI, reads through the standard API, and never touches Northbay's configuration.

Price is set on **active projects**, not seats — Qadira already owns 22 unused seats, so pricing on seats would repeat the mistake that's souring the CFO. An 8-week pilot on 1–2 projects is **AED 85,000** one-time; ongoing service is **AED 6,500 per active project per month**, with a step-down above six concurrent projects. Full detail and the year-two (eight-project) number are in Section 3.

The business case, on estimated Qadira numbers labelled and shown as arithmetic, puts the annual cost of the three problems Layla named at roughly **AED 4.5–5.5M** — against a year-one Groundtruth cost of about AED 397,000. Even discounted 50% for CFO scepticism, that's a **~470–590% year-one ROI** with payback inside the first two quarters of full deployment. That risk-adjusted number, not the headline one, is what goes in front of the CFO. Full workings are in Section 4.

The sell is sequenced Northbay, then Layla, then the CFO — never out of order — with a specific, checkable claim (the stuck tower on floors 7–11) as the credibility hinge for everything else. Section 5 sets out what each meeting looks like, what Northbay gets, and what kills the deal.

SECTION 01

The Offering

1.1 What It Is

Four modules, each answering one thing Layla asked for by name, delivered as a weekly digest and exception alerts rather than a dashboard she has to remember to open:

- **Pipeline Reality Score** — a re-forecast of what's actually going to close this quarter, built from deal velocity, activity recency and historical stage-to-close conversion, shown next to the CRM's own stage-based forecast so the gap is visible rather than asserted.
- **Deal Risk Radar** — a short, ranked list of deals losing momentum right now, each with the specific signal that flagged it (stalled stage age, dropped activity, a missed milestone), refreshed weekly.
- **Broker Net-Lift Ranking** — of the ~90 agencies, which bring buyers Qadira would not otherwise have reached, versus which intercept demand that would have walked in anyway, measured against the direct-channel baseline rather than raw volume.
- **Unit Velocity Map** — days-on-market by block, floor and unit type, so a pocket like floors 7–11 sitting unsold for five months surfaces on its own, with a candidate cause attached where the data supports one.

Delivered Monday morning as a one-page email/PDF — the channel the offer is priced and committed against — with a WhatsApp Business message running alongside it once Meta's template approval is in place. That approval, and WhatsApp's per-conversation cost, is a real dependency we don't control, not a guaranteed day-one channel, so it's presented as the faster option Qadira is already comfortable with, not the only one. To be clear, this is a one-way message Groundtruth sends to Layla — it reads nothing from, and has no connection to, the internal operations

WhatsApp group her team already runs on, which stays explicitly out of scope (3.4). Exception alerts run through the week if a deal or a unit crosses a threshold. The detail behind each headline is available in Salesforce for anyone who wants to click through, but nothing in the sales team's day-to-day changes.

1.2 What It Explicitly Is Not

- Not a second AI product stacked on the add-ons Qadira already owns and doesn't use. Groundtruth doesn't sit inside Salesforce's UI or compete for the adoption Sandeep is worried about; it's an outboard layer that reads via API and pushes out.
- Not a data-entry programme. No new fields, no new required steps. Everything derives from data already captured, plus timestamps Salesforce records automatically.
- Not a Northbay replacement. Northbay keeps the org, the customisation and the rate-card support; Groundtruth reads from the org and never touches configuration Northbay is responsible for.
- Not a general BI platform or a data-warehouse project. Scoped to the four modules tied directly to the four things Layla named — nothing broader.

1.3 Why This Shape

Layla's ask carries a hard constraint (no more data entry) and a hard requirement (told, not asked to look) — both point away from a dashboard and toward a small number of derived signals, pushed. Sandeep's constraints — don't cannibalise the rate card, don't repeat the AI add-on mistake — point toward a narrow product that sits beside Salesforce, with Northbay holding a continuing commercial role rather than being routed around. The CFO's likely objection, unstated but predictable, points toward proving cost avoidance in his numbers before asking him to sign anything ongoing. Section 3 and 4 are built around that last constraint specifically.

1.4 How the Signals Are Calibrated

Two of the four modules carry a calculation question worth answering directly rather than leaving implicit.

Pipeline Reality Score — calibrated against the last three to four closed quarters of data Salesforce already captures automatically: stage-change timestamps, activity-log dates, email and call metadata. None of it is a field the sales team fills in by hand. The pilot's first output is a backtest against those closed quarters, so Layla can see whether the model would have called the last miss correctly before it's asked to call the next one live — which is also the answer to the fair worry that "no more data entry" and "a working forecast model" pull in different directions. They don't, because the signal comes from timestamps the system logs on its own.

Broker Net-Lift Ranking — measured by comparing each broker-sourced buyer's pre-contact activity trail — did they visit a sales centre, view a listing directly, or engage a Qadira channel before the broker's first logged touch — and buyer profile (unit type, price band, project) against the equivalent profile of direct-channel buyers in the same period. A broker whose buyers consistently show no prior Qadira contact and a profile unlike the walk-in baseline scores high lift; one whose buyers repeatedly match people already in Qadira's own pipeline scores low. It's a comparison against a stated baseline, not a black-box score, and every ranking in the digest carries the specific comparison that produced it — the same discipline as the named reason behind every Deal Risk Radar flag.

1.5 Guardrails: Flagged for Review, Not Auto-Acted

Every output is a flag for a person to weigh, not an instruction to execute. This matters most for the broker ranking and the deal risk radar, because acting on a wrong flag costs more than the model being wrong once: cutting a broker relationship on a bad call costs Qadira one of ninety active partners, and overriding a salesperson's read of a deal on a bad call costs trust in the whole system, not just that call. Every flag carries a stated confidence band, and the rule we'd propose to Layla directly is that broker and deal-risk flags go to a review conversation in week one, not an automatic commission or pipeline change, until the model has a validated season — roughly one full quarter — of calls behind it on a given project.

SECTION 02

The Offering, Made Real

A description of a product is not a product, so the pack includes a working prototype (qadira_prototype.html) built on invented but internally consistent Qadira data across five exhibits:

1. Pipeline Reality — the Q3 CRM forecast (AED 61.4M) against the Groundtruth reality score (AED 43.9M), broken out by project with the specific driver of each gap.
2. Deal Risk Radar — four live deals, ranked, each with the exact signal that triggered the flag rather than a generic risk label.
3. Broker Net-Lift — a ranking of four agencies by incremental lift rather than volume, with the reasoning shown.
4. Unit Velocity Map — a floor-by-block heat map of Palm Vista Towers, Tower B, surfacing the floors 7–11 problem and a specific, checkable candidate cause.
5. The Monday Digest — the actual WhatsApp message Layla receives, which is the real product; the other four exhibits are what sits behind it for anyone who clicks through.

The prototype is deliberately built against Layla's own example (the stuck tower) so the first thing she sees in the room is a specific answer to something she already knows is a problem, not a generic capability demo.

SECTION 03

Pricing

3.1 Methodology

Price on the unit that drives both cost-to-serve and value: the active project, not the Salesforce seat. Qadira already owns 22 unused seats and a dormant AI add-on; pricing on seats repeats the exact mistake Sandeep flagged before we're in the room. Each active project is its own pipeline and inventory pool with its own calibration — unit mix, price band and broker footprint differ project to project — so cost to serve scales with project count, and that's also the variable the brief itself asks about directly (four projects going to eight).

Two parts: a one-time Discovery & Build fee per pilot, and a recurring per-project subscription.

3.2 The Pilot

Scope: 1–2 of the four active projects. **Recommended: Palm Vista Towers**, because the floors 7–11 problem gives the risk radar and unit map a specific, checkable answer inside the pilot window rather than an abstract one.

Duration: 8 weeks. Price: **AED 85,000 (~US\$23,150)**, one-time, covering both pilot projects — all four modules, eight weekly digests, and one working session with Layla's team to check the Pipeline Reality Score against what actually closed. No subscription fee during the pilot.

3.3 What Follows the Pilot

Scenario	Basis	Monthly	Annualised
Today (4 projects)	AED 6,500 / project / month	AED 26,000	~AED 312,000
Year two (8 projects)	Volume step-down: 15% off from project 7	~AED 44,000–46,000	~AED 540,000

Headcount does not move the price — the brief's "eight projects and forty more people" scenario is priced entirely on the project count, because that's what drives the calibration work, not the seat count. Each new project beyond the pilot carries a one-time AED 15,000 calibration-only fee (no new integration work) before its subscription starts.

3.4 What's Included / Not Included

Included: all four modules per active project, weekly digest, exception alerts, a monthly 30-minute review call, and recalibration as data accumulates.

Not included: any change to Salesforce configuration or objects (Northbay's rate card); data sources outside Salesforce — specifically, reading or mining Qadira's internal operations WhatsApp group, which is a different thing from the one-way WhatsApp digest Groundtruth sends to Layla (1.1) and stays explicitly out of scope as a separately priced add-on, because it's a different kind of build; training beyond the one included session; a materially different asset class such as villas, priced as a small calibration add-on rather than a rebuild.

3.5 Resilience to a Northbay Org Upgrade

Groundtruth reads through the standard Salesforce API against stable object and field names, not through UI customisation or managed-package hooks a Northbay release could break. The main real risk is a field rename or page-layout change; the subscription includes a compatibility check after any Northbay release flagged to us in advance, at no extra charge — which is also the operational reason Northbay needs to stay in the loop rather than be routed around.

3.6 Northbay's Commercial Position

A 15% referral share on the Discovery & Build fee and on the first twelve months of subscription revenue for any client Northbay introduces, paid to Northbay on top of our price, not deducted from Qadira's. Any org-side change needed to support Groundtruth — exposing a field, adjusting a permission set — is scoped and billed by Northbay under their existing rate card; we do not fold configuration work into our price or touch it ourselves. That keeps the two engagements financially separate and keeps Northbay as the party holding the account relationship, which is the actual reason they introduce us to the other four developer clients Sandeep mentioned, rather than treating this as a one-off.

Delivery capacity for those other four clients is built around the same calibration method in 1.4 rather than a bespoke build each time — discovery is templated, which is what keeps the AED 15,000 per-project calibration fee (3.3) realistic at volume rather than a number that only works once.

3.7 Data Handling, Residency and Cross-Client Confidentiality

Groundtruth connects through a named, read-only integration user scoped to the objects the four modules need — it does not read email content, the operations WhatsApp group, or anything outside Salesforce. Model outputs (the digest and the four exhibits) belong to Qadira; the underlying data is never pooled, benchmarked or shared in any form, aggregated or otherwise, with any other client, including the other developer clients Northbay may introduce. A cross-client benchmarked view, if Qadira and Northbay wanted one in future, would be a separate, explicitly opted-in product, never a default. Data residency and retention — where models and outputs are hosted, and what happens to Qadira's data and access on contract end — are confirmed in Discovery against whatever PDPL and lender or board reporting obligations Qadira already carries. This is a term to close before the pilot starts, not after.

SECTION 04

The Business Case for Qadira

We were not given Qadira's financials, licensing detail or commission structure. Every figure below is an estimate, marked [EST], built from the details in the brief and typical UAE off-plan ranges, with the arithmetic shown so it can be replaced with Qadira's real numbers in the first working session.

4.1 Assumptions Used

- [EST] Average unit price: AED 1,200,000 (~US\$327,000), a mid-market off-plan blend across Dubai/Sharjah.
- [EST] Annual unit sales rate: 480–600 units/year across the 2,400-unit, four-project portfolio.
- [EST] Broker-sourced share of sales: 35–40% of units, consistent with a ~90-agency channel alongside two walk-in centres and portal leads.
- [EST] Standard broker commission: 2% of gross sale value.
- [EST] Cost of capital / construction financing: 8% per annum, applied to delayed revenue recognition.
- [EST] Carrying cost per unsold unit per month: AED 3,500–5,000 (financing, marketing and opportunity cost).

4.2 Cost of the Forecasting Miss

At AED 150–180M in quarterly sales (480–600 units/year ÷ 4, at AED 1.2M average), a miss "by a third" against the board-reported forecast implies roughly AED 50–60M of expected revenue that did not land in-quarter.

EXHIBIT 4.2 — FINANCING COST OF A ONE-QUARTER DELAY	
Quarterly sales run-rate [EST]	AED 150M – 180M
Revenue delayed by the miss [EST]	AED 50M – 60M
Cost of capital, annualised [EST]	8%
<i>Formula</i>	<i>Delayed revenue × cost of capital ÷ 4 quarters</i>
Financing cost, this quarter alone	≈ AED 1.1M

The cost that doesn't show up as a line item — credibility with the board and, if Qadira carries development finance, with lenders reading the same board pack — is the one Layla is actually describing; AED 1.1M is offered as a floor, not the real cost.

4.3 Cost of Unpriced Broker Spend

At an assumed 40% broker share, roughly 220 units a year run through the channel.

EXHIBIT 4.3 — UNPRICED BROKER COMMISSION	
Broker-sourced units, annual [EST]	≈ 220 units
Average unit price [EST]	AED 1.2M
Standard commission rate [EST]	2.0%
<i>Formula</i>	<i>Units × price × commission rate</i>
Total annual commission paid	≈ AED 5.3M
Estimated no-lift share [EST]	20%
Wasted commission spend, annual	≈ AED 1.06M

The broker module doesn't need to prove the exact figure to be worth building — it turns "I suspect" into a ranked list Layla can act on at the next commission review.

4.4 Cost of Stuck Inventory

Using the floors 7–11 example as the unit of measurement:

EXHIBIT 4.4 — CARRYING COST, TOWER B, FLOORS 7–11	
Unsold units in the affected block [EST]	≈ 40 units
Months unsold [EST]	5
Carrying cost per unit, per month [EST]	AED 4,000
<i>Formula</i>	$Units \times months \times carrying\ cost\ per\ unit$
Cost tied up, this block alone	≈ AED 800,000
Portfolio run-rate (4 projects, 1 pocket each)	≈ AED 2.5M – 3.5M / yr

Assuming one comparable pocket of stuck inventory per project at any given time — a conservative read across four active projects — this is the run-rate an earlier signal could meaningfully reduce, even without eliminating it.

4.5 Reading the Numbers Together

Total plausible annual cost across the three problems Layla named: roughly AED 4.5–5.5M/year [EST — sum of 4.2 annualised, 4.3 and 4.4], against a year-one Groundtruth cost of about AED 397,000 (AED 85,000 pilot plus AED 312,000 for a full year across four projects, noting the pilot fee doesn't repeat). That's the frame for the CFO — cost avoidance against costs he already carries — before any discounting. Section 4.6 turns that into ROI and payback, and applies a discount, because a CFO who already regrets one vendor purchase should not be handed an unadjusted number and asked to trust it.

4.6 Return on Investment and Payback

Two figures, deliberately: the unadjusted case, and a risk-adjusted case that assumes the estimates in 4.1–4.4 are only half right. Cost-avoidance claims are the easiest kind of ROI to oversell, and the CFO's own history with this vendor relationship is a purchase that didn't earn its adoption — so the risk-adjusted figure, not the headline one, is what should go in front of him first.

	Unadjusted	Risk-adjusted (50% haircut)
Estimated annual benefit	AED 4.5M – 5.5M	AED 2.25M – 2.75M
Year-one Groundtruth cost	AED 397,000	AED 397,000
Net benefit, year one	AED 4.1M – 5.1M	AED 1.85M – 2.35M
Return multiple (benefit ÷ cost)	11.3× – 13.9×	5.7× – 6.9×
ROI (net ÷ cost)	~1,030% – 1,290%	~470% – 590%

EXHIBIT 4.6 — PAYBACK, RISK-ADJUSTED	
Year-one Groundtruth cost	AED 397,000
Risk-adjusted annual benefit	AED 2.25M – 2.75M

<i>Formula</i>	$Cost \div (annual\ benefit \div 12)$
Payback, at full run-rate	$\approx 6 - 9\text{ weeks}$

Payback doesn't start on day one, though — the pilot spends its first eight weeks proving the model, not yet running at full run-rate, and some benefit categories act faster than others: the unit-velocity fix on floors 7–11 is checkable within the pilot itself, while a broker-commission decision only saves money at the next commission cycle, and the forecast-accuracy claim only earns full trust after it's been right for a quarter. The honest payback claim is not "seven weeks" — it's that even on the risk-adjusted, half-credit numbers, the programme should have paid for itself within the first two quarters of full deployment, with the unit-velocity and deal-risk benefits landing earliest and the forecast and broker benefits landing by the second quarter.

Three-Year View, Tied to Section 3 Pricing

Using the same linear scaling assumption as the pricing model — benefit scales with project count in roughly the same proportion as cost — and holding the 50% haircut constant:

Year	Projects	Annual cost	Risk-adj. benefit	Risk-adj. ROI
Year 1	4 (incl. pilot)	~AED 397,000	~AED 2.5M	~530%
Year 2	6 [illustrative]	~AED 468,000	~AED 3.75M	~700%
Year 3	8	~AED 540,000	~AED 5.0M	~825%

This assumes each additional project carries roughly proportional pipeline, broker and inventory problems to the first four — a simplifying assumption, not a forecast, and one that should probably be conservative in Qadira's favour: a larger portfolio more often has compounding coordination costs (more brokers to track, more towers competing for the same buyer pool), which would push real benefit above this line rather than below it. It's shown to answer the year-two/year-three cost question from the CFO's own numbers, not to promise a specific return.

SECTION 05

How We Sell It

5.1 Sequence

Northbay first, always. Sandeep is the gate — without his introduction there's no meeting with Layla, and going to Layla or the CFO out of sequence burns this deal and the other four he mentioned.

Meeting 1 — Northbay, before Layla. Not a pitch — a commercial conversation. Walk Sandeep through what we're proposing to say to Layla, the referral structure (3.6), and the explicit boundary that we don't touch org configuration. Ask what he needs to be comfortable making the introduction, and adjust before we're in front of the client. What Northbay gets: 15% referral economics on a deal he doesn't have to build or support, a stickier Qadira account, and a reference for the other four developer clients if it goes well.

Meeting 2 — Layla, with Sandeep in the room. Opens with the prototype on her own tower, not a slide deck. We show the four exhibits against invented-but-plausible data, explicit that it's a mock-up we need her real data to correct. The ask at the end is not a signature — it's the introduction to the CFO and agreement on which project(s) to pilot.

Meeting 3 — the CFO, with Layla present. Opens with Section 4's numbers, not the product. We name the AI add-on adoption problem ourselves in the first two minutes rather than let him raise it, and show why this is structurally different: outboard, not inside Salesforce; priced on projects Qadira already tracks, not seats; an 8-week pilot with a defined stop point, not a multi-year commitment.

5.2 What Each Side Asks, and the Honest Answer

- Layla won't ask about risk; she'll ask whether her team has to do anything differently. Answer: no — nothing changes in how the 22 sales staff work.
- Sandeep will ask what happens at the next org upgrade and whether this becomes support debt for his team. Answer: no — compatibility checks are on us (3.5), and we only need read access.
- The CFO will ask what's lost if it doesn't work. Answer: the AED 85,000 pilot fee and eight weeks. No new customisation, no new licences, no multi-year contract, nothing that outlives the upgrade cycle in 3.5 — that's the actual answer, worth saying exactly that plainly in the room.
- The CFO will also ask, in some form, why not just give the existing analyst a sharper brief instead of paying an outside vendor. Answer: capacity, not capability — that analyst is already at full stretch keeping three spreadsheets current by hand, and the behavioural comparisons behind 1.4 (pre-contact activity trails, baseline matching) aren't a spreadsheet-native calculation to begin with.

5.3 What Kills the Deal

- Positioning this as a second AI product, even by accident, in front of the CFO.
- Any suggestion, in Sandeep's hearing, that we'd rather deal with Qadira directly than through Northbay.
- Asking Layla's team to change a single habit in week one — even a small one undoes the one condition she stated outright.
- A pilot that can't produce a checkable, specific answer (the stuck tower) inside eight weeks — it's the credibility hinge for everything else.

5.4 Trust, Not Just Accuracy

The harder adoption risk isn't whether Layla's team changes behaviour — it's what happens the first time the model contradicts someone she trusts. If the Deal Risk Radar flags a deal her best salesperson swears is fine, the model needs to survive that moment, not just be right on average across a quarter. Two things make it survivable rather than a one-shot credibility event: every flag carries the specific reason it fired (1.1, 1.4), so it's a conversation about a fact — no contact in 24 days — rather than a verdict to accept or reject outright; and the guardrail in 1.5 means the first disagreements go to review rather than automatic action, so the model earns trust on calls Layla can check herself rather than ones she has to take on faith.

SECTION 06

Judgement Calls, Recorded

Where the brief was ambiguous, or the two notes pulled in different directions, here is the judgement made and why, rather than leaving it as an unstated assumption:

1. Chose a narrower four-module offering over a broader analytics platform, on the brief's own framing that a narrow, honestly priced offering beats a broad one that can't be costed.
2. Assumed email/PDF as the committed delivery channel and WhatsApp as a faster add-on rather than the reverse, since WhatsApp Business template approval is outside our control — flagged as a dependency to close in week one, not a guaranteed day-one channel (1.1).
3. Priced per active project rather than per seat or per transaction, because it's the variable Qadira already tracks and the one the brief's own "four to eight projects" question implies should drive cost.
4. Read Sandeep's separation of "his rate card" from what we propose as a request for a clean boundary, not a merged offer — so the commercial structure is a referral relationship, not a joint bundle.

5. Where Layla wants to be told things and Sandeep is wary of anything AI-labelled, resolved by delivering statistical/behavioural signals as a push digest rather than branding it a generative-AI feature — same outcome for Layla, without the trigger word for the CFO.
6. Treated the CFO as the hardest audience by default, since he sent no note, and built Section 4 to be defensible without him having said anything at all, rather than assuming he'd be the easiest sign-off.

SECTION 07

Anticipated Questions

From Section 4 of the brief, illustrative rather than exhaustive — answered with a reference back to where the full reasoning sits.

- 1. What does this do that a Salesforce report does not?** A report is pulled and reflects the stage field; Groundtruth is pushed and reflects derived behavioural signals — the whole point of Layla's note is that she doesn't want to have to go and look, and doesn't believe the stage field. See 1.1 and 1.3.
- 2. We already paid for the AI add-ons. Why pay you as well?** Different shape of product entirely — outboard, not inside the CRM competing for the same adoption; narrow and tied to four named problems, not a general capability nobody asked to use. See 1.2.
- 3. When your number and the CRM number disagree, which goes in the board pack?** Groundtruth's — because the CRM's stage-based number is exactly the one Layla said she didn't believe — but shown alongside the CRM figure, not instead of it, so the gap itself becomes a visible trend rather than a disputed replacement. See 2 and Exhibit 01.
- 4. What does this cost in year two, at eight projects and forty more people?** Roughly AED 44,000–46,000/month; headcount doesn't move the price, project count does. See 3.3.
- 5. What's the smallest version we could buy, and what would it prove?** The pilot itself — 1–2 projects, eight weeks — proves whether the stuck-tower flag matches what Layla already knows and whether the forecast gap is real and specific, before any ongoing commitment. See 3.2.
- 6. What happens when Northbay next upgrades the org and your layer breaks?** Compatibility checks are included and run on any flagged release, at no extra charge, because the integration reads stable API objects rather than UI customisation. See 3.5.
- 7. Why should Northbay introduce you rather than build this themselves?** Northbay is an implementation and support business on a rate card; a statistical signal product is a different discipline and a distraction from what they bill for. The referral economics (3.6) give them the upside without the build cost or the support burden.
- 8. If it doesn't work, what have we lost?** AED 85,000 and eight weeks. No customisation, no new licences, no multi-year contract, nothing that outlives an org upgrade cycle. See 5.2.
- 9. Why not just give your analyst a sharper brief, or build this in Excel?** Capacity, not capability — the analyst is already at full stretch maintaining three spreadsheets by hand, and the behavioural comparisons behind the broker and pipeline modules aren't a spreadsheet-native calculation to begin with. See 1.4 and 5.2.
- 10. Does our data get pooled or benchmarked against your other developer clients?** No, not by default and not without an explicit separate opt-in — each client's data and model are isolated. See 3.7.